



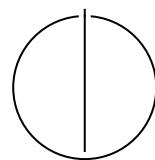
SCHOOL OF COMPUTATION,
INFORMATION AND TECHNOLOGY —
INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Master's Thesis in Informatics

Constructing Linear Types in Isabelle/HOL

Felix Krayer





SCHOOL OF COMPUTATION,
INFORMATION AND TECHNOLOGY —
INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Master's Thesis in Informatics

Constructing Linear Types in Isabelle/HOL

**Konstruktion linearer Typen in
Isabelle/HOL**

Author: Felix Krayer
Examiner: Florian Bruse
Supervisor: Dmitriy Traytel, Tobias Nipkow
Submission Date: 13-11-2025

I confirm that this master's thesis is my own work and I have documented all sources and material used.

Munich, 13-11-2025

Felix Krayer

Acknowledgments

Abstract

Contents

Acknowledgments	iii
Abstract	iv
1 Introduction	1
2 Background	3
2.1 Bounded Natural Functors (BNFs)	3
2.1.1 Map function and functors	4
2.1.2 Set functions and naturality	5
2.1.3 Bound and boundedness	6
2.1.4 Relator and shapes	6
2.1.5 Additional BNF-axioms	7
2.1.6 Non-emptiness witnesses	8
2.1.7 BNF examples	8
2.2 Syntaxes with bindings	9
2.3 Map-Restricted Bounded Natural Functors (MRBNFs)	9
2.3.1 MRBNF axioms	11
2.3.2 Datatypes with bindings	12
3 Linearizing MRBNFs	13
3.1 Linearizing MRBNFs	13
3.1.1 Non-repetitiveness	13
3.1.2 Conditions for linearization	14
3.1.3 Intermediate lemmas	15
3.1.4 Defining the subtype and its constants	20
3.1.5 Proving the MRBNF axioms	21
3.1.6 Lifting Witnesses	22
3.1.7 Preservation of strength	22
3.2 Implementing the <code>linearize_mrbnf</code> command	23
4 Examples	26
4.1 POPLmark Challenge: Pattern	26

Contents

Abbreviations	29
List of Figures	30
List of Tables	31
Bibliography	32

1 Introduction

Isabelle/HOL provides a fleshed out system for constructing datatypes and codatatypes through the **datatype** command that is built on the theory of Bounded Natural Functors (BNFs).

This allows to express and thus reason about many different kinds of structured concepts.

How about expressing lambda calculus terms? As an example we consider lambda terms with parallel let bindings, where a term is either a variable, a lambda abstraction, an application or a parallel let binding. Concretely, we want to represent the following syntax, where x are variable names.

$$t := x \mid (\lambda x. t) \mid t_1 t_2 \mid \text{let } x_1 = t_1 \text{ and } \dots \text{ and } x_n = t_n \text{ in } t_{n+1}$$

We can express the structure of these terms with a regular Isabelle datatype as follows, where α is the type of variable names:

$$\begin{aligned} \text{datatype } \alpha \text{ lterm} = & \text{Var } \alpha \mid \text{Abs } \alpha \ " \alpha \text{ lterm}" \mid \text{App } " \alpha \text{ lterm}" \ " \alpha \text{ lterm}" \\ & \mid \text{Let } "(\alpha \times \alpha \text{ lterm}) \text{ list}" \ " \alpha \text{ lterm}" \end{aligned}$$

However, this representation is not ideal. For example, it does not consider α -equivalence of lambda terms, i.e., the notion that $(\lambda x. x z)$ and $(\lambda y. y z)$ are semantically equivalent terms.

Furthermore, this datatype represents invalid terms. This makes working with the type very tedious, since it calls for a separate definition of and checks for validity. For example, α lterm contains "Let $[(\text{Var } x, \text{Var } y), (\text{Var } x, \text{Var } z)]$ ($\text{Var } x$) as a valid term. This would be the representation of "let $x = y$ and $x = z$ in x " which is not a valid type, since a variable cannot be bound to different terms in the same parallel let binding. This can in theory be solved by defining a subtype of $(\alpha \times \beta)$ list that only contains lists of pairs, where the first elements of each pair are unique in the list. However, this subtype is not a BNF right after that and thus it cannot be used in the recursive definition.

An important step towards solving these issues has been taken by Blanchette et al. [1] in their implementation of a definitional package for binding-aware datatypes in Isabelle/HOL. They provide a new **binder_datatype** command that allows for the

definition of datatypes that do not distinguish between α -equivalent terms and provide functions for renaming and obtaining the free variables of a term. Furthermore these binding-aware datatypes (or short *binder datatypes*) provide propositions similar to those provided by a regular datatype, e.g., a principle for structural induction. This package is built on the theory of Map-Restricted Bounded Natural Functors (MRBNFs), a generalization of BNFs.

It is possible to carve out subtypes with Isabelle's **typedef** command (and quotient?), however thereby the BNF-Properties are lost.

While this solves some issues, it relies on manually linearizing types, i.e. creating types that are non-repetitive. These are not closed under ... and thus they cannot be structurally defined. Instead a non-linear type is built and then they are carved out. This was a manual process but now we automate it.

- Datatypes in Isabelle/HOL are built on BNFs (defined in [2])
 - Difficult represent lambda terms (alpha equivalence, capture-avoiding substitution)
- => binding-aware datatypes [3]
- Structure of the Thesis

2 Background

This Chapter serves to introduce Bounded Natural Functors (BNFs) and their generalization to Map-Restricted Bounded Natural Functors (MRBNFs). Note, that when we use the notion "element" of a n -ary type constructor F , we are always talking about a term of type $(\alpha^n) F$. In contrast to that, we call the α_i elements that make up a F element "atoms". The structure of the F element dictates how and where the atoms occur in it.

2.1 Bounded Natural Functors (BNFs)

As described in Chapter 1, BNFs are essential for constructing datatypes and codatatypes in Isabelle/HOL. Especially for defining a datatype with recursion it is required that the type constructor used in that recursion is registered as a BNF, i.e., it fulfills the BNF-axioms. For example the following **datatype** command only succeeds if α list and $(\alpha \times \beta)$, are BNFs.

```
datatype α ex = A "((α × α) ex) list"
```

Since BNFs are closed under composition and fixpoints, the resulting datatype (here α ex) can be automatically registered as a BNF as well.

We write type variables as greek letters (α, β, \dots) in this thesis. However, in the Isabelle proof assistant type variables are written with a `''` in front of a name, e.g., 'a list. To copy our examples to Isabelle, one has to replace these greek letters with `''` variables. Alternatively, a `''` can be prepended to the greek letters, since for example `'α` is a valid type variable in Isabelle.

The type variables of a BNF are divided into two groups: *live* and *dead* variables or *lives* and *deads*. Recursive occurrences may appear only in positions corresponding to live variables when defining a new datatype. Dead variables do not allow for this. We take the function type $\alpha \Rightarrow \beta$ as an example. It's first type argument α is dead, while the second one β is live. Thus, of the following the first command succeeds while the second one fails:

```
datatype α success = S1 | S2 "α ⇒ α success"  
datatype α fail = F1 | F2 "α fail ⇒ α"
```

The reason for this failure is, that certain properties have to hold for a BNF with regard to its live variables. These properties are necessary for the internal construction of a newly specified type. A variable is dead when it has to be omitted such that these axioms hold, which is the case for the first type variable of the function type.

These properties - or "BNF-axioms" as we call them - are formalized in terms of constants that characterize a BNF. For a BNF F with l live variables these are one $l + 1$ -ary map function, l set functions, a bound and a $l + 2$ -ary relator. In the following Subsections Subsection 2.1.1 - Subsection 2.1.4 we define these constants and motivate the BNF-axioms that we formalize in Figure 2.1. We use the notation $f^l = f_1 \dots f_l$ for the arguments of the mapper and similarly $R^l = R_1 \dots R_l$ for the relator. Note that we extend this notation to binary (infix and prefix) operations op on functions and relations as follows: $(R^l \text{ op } Q^l) = (R_1 \text{ op } Q_1) \dots (R_l \text{ op } Q_l)$ and $(\text{op } R^l \text{ } Q^l) = (\text{op } R_1 \text{ } Q_1) \dots (\text{op } R_l \text{ } Q_l)$.

- (MAP_ID) $\text{map}_F \text{id}^l x = x$
- (MAP_COMP) $\text{map}_F g^l (\text{map}_F f^l x) = \text{map}_F (g \circ f)^l x$
- (MAP_CONG) $(\forall i. \forall z \in \text{set}_{F,i} x. f_i z = g_i z) \implies \text{map}_F f^l x = \text{map}_F g^l x$
- (SET_MAP) $\forall i. \text{set}_{F,i} (\text{map}_F f^l x) = f_i ` \text{set}_{F,i} x$
- (BD) infinite $\text{bd}_F \wedge$ regular $\text{bd}_F \wedge$ cardinal_order bd_F
- (SET_BD) $\forall i. |\text{set}_{F,i} x| <_o \text{bd}_F$
- (REL_COMPP) $(\text{rel}_F R^l \bullet \text{rel}_F Q^l) x y \implies \text{rel}_F (R \bullet Q)^l x y$
- (IN_REL) $\text{rel}_F R^l x y = \exists z. (\forall i. \text{set}_{F,i} z \subseteq \{(a,b). R_i a b\}) \wedge \text{map}_F \text{fst}^l z = x \wedge \text{map}_F \text{snd}^l z = y$

where ` is the image function on sets, \bullet is the composition of relations and $<_o$ is the less than relation on cardinals

Figure 2.1: The BNF axioms

2.1.1 Map function and functors

The map function or *mapper* takes one function for each live of F as arguments as well as one F element. The domain types of these functions are the lives of F . These functions are applied to the atoms of an element. The result is a new element of type F , where the original type variables are replaced by the range types of the mapped

functions. Taking the α list type as an example, a BNF with one live α , the mapper has the type $\text{map}_{\text{list}} :: (\alpha \Rightarrow \alpha') \Rightarrow \alpha \text{ list} \Rightarrow \alpha' \text{ list}$.

To make F with its mapper a *functor* on the universe of all types, the mapper has to fulfill two axioms [2]. First, mapping the id function on all lives over an element should leave it unchanged, which is formalized in **MAP_ID** (Fig. 2.1). The second property **MAP_COMP** (Fig. 2.1) is concerned with mapping composed functions and reads as follows: Mapping two lists of functions over an element, e.g., first $f_1 \dots f_l$ and then $g_1 \dots g_l$, should produce the same result as mapping the pair-wise composed functions $(g_1 \circ f_1) \dots (g_l \circ f_l)$ over it once. A type constructor F with a map function map_F fulfilling these two properties is considered a functor.

2.1.2 Set functions and naturality

A set function or *setter* is defined for each of the l live variables. Applied to an F -element, the i -th setter returns the set of all atoms that are part of the element and correspond to the i -th live. For example, the setter of the list type returns the set of elements in the list. We note here that when we write i as an index, we assume it to be in the range $1 \leq i \leq l$.

The set functions together with the mapper give rise to another property. We want the setters $\text{set}_{F,i}$ to be natural transformations from F and map_F to the set and image function. Thus, they should fulfill the **SET_MAP** axiom (Fig. 2.1). It states that taking the i -th set of an F after mapping $f_1 \dots f_l$ to it, results in the same set as if i -th set was taken from the original F before the image of f_i was applied to it. Figure 2.2 shows a visualization of this axiom and reads as follows: Starting from an F element first applying the setter and then mapping a function (path through the top right) results in the same as first mapping the function and then applying the setter (path through the bottom left).

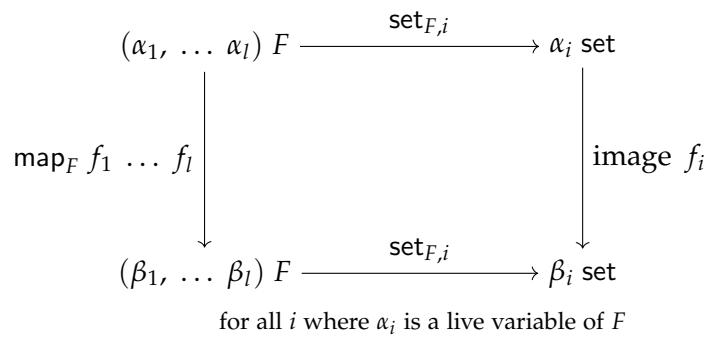


Figure 2.2: $\text{set}_{F,i}$ as a natural transformation

This axiom alone would be fulfilled by declaring every setter as the constant function returning the empty set. To solve this, the `MAP_CONG` axiom (Fig. 2.1) acts as a completeness property on the setter. It states that if two (lists of) functions f^l and g^l are equal when applied to the corresponding sets of all atoms of an F element (obtained through the setters), then mapping these two lists of functions over the F element each produces the same result. When this property holds, we can be sure, that the mapper only depends on how the functions f^l behave on the atoms of the F element. At the same time this axiom ensures that the setters actually return the complete set of atoms of the F element.

2.1.3 Bound and boundedness

Lastly, the BNF must be equipped with an infinite cardinal as a bound. This bound may depend on the cardinalities of the dead variables, but not on the of the live variables. In Isabelle/HOL cardinals are implemented as minimal wellorders with respect to isomorphisms [4]. For example `natLeq`, the cardinal that originates from the \leq order on natural numbers, is equivalent to the smallest infinite cardinal \aleph_0 . While details about this implementation are certainly interesting, we will not focus on these details in this thesis and refer to cardinals in terms of their \aleph -notation.

Besides being a cardinal order, the bound is required to be infinite, i.e., at least \aleph_0 with respect to the cardinal order \leq_o , and regular. Regularity means that an infinite cardinal κ is stable under union, i.e., the union of any set of sets that are of smaller cardinality than κ also has smaller cardinality than κ . Note that the set of sets must also be of smaller cardinality than κ :

$$(\bigwedge_{i \in I} |S_i| <_o \kappa) \wedge |I| <_o \kappa \implies |\bigcup_{i \in I} S_i| <_o \kappa$$

The bound is used in the `SET_BD` axiom (Fig. 2.1) to ensure that the sets obtained by the setters are bounded. This ensures that the branching of a recursively defined datatype is also bounded and so is the resulting type F as well.

2.1.4 Relator and shapes

The relator is used to build a relation on F by relating the atoms of an F element. It takes one relation for each live, that relates the corresponding type variables of the two F s that are to be related. As an example we give the type and definition of the relator

for the product and list type as follows, where $x!i$ refers to the atom of list x at index i :

$$\begin{aligned}\text{rel}_{\text{prod}} &:: (\alpha \Rightarrow \alpha' \Rightarrow \text{bool}) \Rightarrow (\beta \Rightarrow \beta' \Rightarrow \text{bool}) \Rightarrow (\alpha \times \beta) \Rightarrow (\alpha' \times \beta') \Rightarrow \text{bool} \\ \text{rel}_{\text{prod}} R Q x y &:= R (\text{fst } x) (\text{fst } y) \wedge Q (\text{snd } x) (\text{snd } y) \\ \text{rel}_{\text{list}} &:: (\alpha \Rightarrow \alpha' \Rightarrow \text{bool}) \Rightarrow \alpha \text{ list} \Rightarrow \alpha' \text{ list} \Rightarrow \text{bool} \\ \text{rel}_{\text{list}} R x y &:= \text{length } x = \text{length } y \wedge (\forall i \leq \text{length } x. R (x!i) (y!i))\end{aligned}$$

Considering the list type again, we make an interesting observation: There are some α lists x and y that the relator cannot relate, no matter which α relation is chosen. The relator on lists is defined index-wise, i.e., the α relation must relate the elements of both lists for each index. Consequently lists of different length cannot be positively related. We think of the length of a list as its *shape*. We can generalize this idea of shape to an arbitrary type constructor F . The shape of an F element is defined by the way it is constructed and the relator can only ever relate those that have the same or equivalent shape, i.e., it will always evaluate to *false*, when two elements of different shape are given, regardless of the relations given to the relator. We can think of an element of type F as a container that has a certain *shape* with slots for *atoms*. These atoms are elements of the type constructor's type arguments.

2.1.5 Additional BNF-axioms

Additionally to the axioms we already motivated in the previous subsections (`MAP_ID` and `MAP_COMP` for the functoriality of F , `SET_MAP` and `MAP_CONG` to ensure that the setters are natural transformations and the boundedness of the setters `SET_BD`), we have three additional ones.

The axiom `BD` (Fig. 2.1) just ensures that the bound bd_F is a suitable cardinal, i.e., a regular and infinite one.

The relator is required to be distributive, i.e., it should fulfill $(\text{rel}_F R^l \bullet \text{rel}_F Q^l) x y = \text{rel}_F (R \bullet Q)^l x y$. We note here, that for showing that a type constructor is a BNF, it is only necessary to prove the implication stated in `REL_COMPP` (Fig. 2.1). The other direction of the implication follows automatically.

Lastly, `IN_REL` (Fig. 2.1) characterizes the relator. It is the most abstract and complex axiom but we want to give an intuition about it here: The idea is that two elements x and y of the type $(\alpha^l) F$ are related through a relation R iff there exists a z that acts as a "zipped" version of x and y . The atoms of this z are R_i -related pairs of the atoms of x and y , where the first position in the pair corresponds to x and the second one to y . This axiom (`IN_REL`) together with the previous one (`IN_REL`) amounts to *weak pullback preservation* of the BNF.

2.1.6 Non-emptiness witnesses

BNF carry non-emptiness witnesses as proof that the type contains at least one element. Witnesses may depend on a subset of the BNF's live variables. For example a witness of $(\alpha_1, \dots, \alpha_l) F$ that depends on the first and last type variable of F , this witness has the type $\text{wit}_F :: \alpha_1 \Rightarrow \alpha_l \Rightarrow (\alpha_1, \dots, \alpha_l) F$. It denotes that given witnesses for the types α_1 and α_l , a witness for F can be constructed. A deeper insight into the theory and usage of witnesses for datatypes in Isabelle/HOL is given by Blanchette et al. [5].

Witnesses have to fulfill the following properties: For all type variables α_i the witness depends on, the witness may only contain the α_i elements w_i , that were given to the witness as arguments, i.e., $\text{SET}_{F,i}$ applied to the witness evaluates to the singleton $\{w_i\}$. Furthermore, the witness must not contain any elements of the live type variables α_j , the witness does not depend on, i.e., $\text{SET}_{F,j}$ must be empty. We formalize these properties in the following where \bar{w} denotes the arguments that the witness depends on.

$$(\text{WITS}) \quad \forall i. \text{set}_{F,i} (\text{wit}_F \bar{w}) = (\text{if } \text{wit}_F \text{ depends on } \alpha_i \text{ then } \{w_i\} \text{ else } \emptyset)$$

Multiple witnesses can exist for a given type constructor. For example, we can give two witnesses for α list: $[]$ of type α list and $(\lambda a. [a])$ of type $\alpha \Rightarrow \alpha$ list. The first of these witnesses does not depend on a type variable, while the second one depends on α . The first witness $[]$ is the more general of the two, since it allows us to give an element of the list type without the need for an element of α . In general, we say a witness $\text{wit}_{F,1}$ *subsumes* $\text{wit}_{F,2}$, when $\text{wit}_{F,1}$ depends on a proper subset of the arguments of $\text{wit}_{F,2}$. Subsuming witnesses are more versatile than subsumed ones and thus we ignore subsumed ones when registering them with a BNF. However, when two witnesses have overlapping dependencies but neither depends on a subset of the other we are interested in both, since none fully subsumes the other.

2.1.7 BNF examples

Further examples of BNFs are the product type (α, β) prod, a binary type constructor with infix notation $\alpha \times \beta$, and the type of finite sets α fset. The latter is interesting for the reason that it is a subtype of the set type, which is not a BNF. By enforcing finiteness for the elements of the type it is possible to give a bound for the set function, fulfilling the `SET_BD` axiom, which is not possible for the unrestricted set type. Since unboundedness is the only reason that the set type is not a BNF, α fset can be shown to be a BNF.

To show, how BNFs can be combined to create new ones, we consider the type constructor (α, β) plist = $(\alpha \times \beta)$ list. We define for it a map function ($\text{map}_{\text{plist}}$) and

two set functions ($\text{set1}_{\text{plist}}$ and $\text{set2}_{\text{plist}}$) as well as a relator $\text{rel}_{\text{plist}} R Q$. We give the exact definitions in terms of the standard map, set and relator functions of the list and product type as follows:

$$\begin{aligned}\text{map}_{\text{plist}} f g &= \text{map}_{\text{list}} (\text{map}_{\text{prod}} f g) \\ \text{set1}_{\text{plist}} xs &= \text{set}_{\text{list}} (\text{map}_{\text{list}} \text{set1}_{\text{prod}} xs) \\ \text{set2}_{\text{plist}} xs &= \text{set}_{\text{list}} (\text{map}_{\text{list}} \text{set2}_{\text{prod}} xs) \\ \text{rel}_{\text{plist}} R Q &= \text{rel}_{\text{list}} (\text{rel}_{\text{prod}} R Q)\end{aligned}$$

To proof that (α, β) plist is a BNF, we have to prove the BNF-axioms for it. Besides the definitions above, we give \aleph_0 as the bound bd_{plist} .

2.2 Syntaxes with bindings

WIP: Considering a polymorphic type that is meant to represent simple λ -terms, where α is the space of variable names. If we want to substitute a free variable x in a term T by a term N , we may run into the following problem: If $T = \lambda y. T'$, we need to ensure that there are no name clashes with y in the new term N before we substitute x by N in T' . This is done by choosing a fresh y' and renaming y to y' in T' .

2.3 Map-Restricted Bounded Natural Functors (MRBNFs)

Type constructors that involve names or bindings often violate the requirements of BNFs. An example of this is the type of labeled lists $\alpha \beta$ alist with α as the type of labels. This is a subtype of the previously defined $\alpha \beta$ plist that describes only lists of pairs that are pairwise distinct on the first elements, i.e., the α atoms. We call it labeled list since one can consider the first atom of each pair a distinct label that identifies the second atom. An alternative but structurally equal interpretation is that of a "key-value" list.

The main issue with this type is that the map function we defined in Subsection 2.1.7 for plist cannot be used for alist right away. The reason is, that it cannot guarantee that the list that results from the map is still distinct on α , i.e., that it is still a member of the type. Thus, in BNF terms the type variable α of (α, β) alist is dead, which is a huge drawback for the versatility of the type. For example, using (α, β) alist in a **datatype** command would kill α also in the resulting datatype effectively disabling any map functions on that type variable. However, by enforcing that only *bijections* are mapped over α , we can ensure that the result of a map on an alist still fulfills the distinctness property of the type. We note here, that there might exist a map function with which

alist could be proven to be a BNF with two live variables. However, it is not obvious how this map function should be defined.

MRBNFs are a generalization of BNFs. Restricting the map function of a functor to *small-support* functions or *small-support bijections* for certain type variables allows us to reason about type constructors in terms of BNF properties, even in cases where this would not be possible otherwise. Concretely, it allows us to include type variables in the mapper that would be considered *dead* from a BNF point of view. We can reason about these variables in terms of the BNF axioms (Fig. 2.1) with a few restrictions and consequently use them in recursive definitions of binder-datatypes under certain conditions. This is explained in more detail in Subsection 2.3.2.

We call type variables that are restricted to small-support functions *free* variables or *frees* and those restricted to small-support bijections *bound* variables or *bounds*. This allows us to define MRBNFs with four types of variables (lives, frees, bounds and deads) as opposed to BNFs which only distinguish between lives and deads. Our example from the beginning of this section, the distinct list α dlist is a MRBNF with α as a bound variable.

A small-support function leaves most arguments unchanged, meaning it acts as the identity function on them. Concretely defined, the cardinality of the set of arguments the function changes must be smaller than the cardinality of the argument type itself:

$$\text{small_supp } f = |\{x :: \alpha. f x \neq x\}| <_o |\Omega_\alpha|$$

where Ω_α is the universe of type α .

For a MRBNF F with l lives, fr frees and b bounds we define $ws = l + fr + b$ as the number of all non-dead type variables. With this, the mapper and setters are expanded to work for the frees and bounds just as they do for lives. Thus, F has ws setters and a mapper with arity $ws + 1$. We note here, that any small-support function acts as the identity function at least for "some" (actually for "most") inputs and thus its domain type and range type are the same. This means that the type variables for frees and bounds are the same for the F argument of the mapper and the result, while the live type variables may be changed through a map.

As before, we use the notation $f^l = f_1 \dots f_l$ for functions and $R^l = R_1 \dots R_l$ for relations on the live variables. Analogously we write v^{fr} and u^b for functions on frees and bounds. Furthermore, we write the arguments of the map function as $f^l v^{fr} u^b$. As an example, the mapper of the type $(\alpha, \beta, \gamma) F$ where α and β are free and γ is bound has the following type:

$$\text{map}_F :: (\alpha \Rightarrow \alpha) \Rightarrow (\beta \Rightarrow \beta) \Rightarrow (\gamma \Rightarrow \gamma) \Rightarrow (\alpha, \beta, \gamma) F \Rightarrow (\alpha, \beta, \gamma) F$$

From now on we assume that the type variables of any MRBNF are ordered *lives* first, followed by *frees* and *bounds*, and *deads* at the end. This simplifies many definitions

and arguments we make about MRBNFs, however these are all easily generalized to an arbitrary ordering of type variables. For example, the argument order of the mapper might be different, as the lives, bounds and frees do not have to be separated, but can be interlaced in some order. In concrete examples we may use MRBNFs that are explicitly defined in terms of primitive types like list or prod. For these examples we may define a different order for live, free and bound type variables.

We keep the original relator that only relates live variables with given relations and relates the free and bound variables with equality. Thinking in our model of F elements being shapes with atoms in slots, the regular relator rel_F requires the frees and bounds in each slot to be the same for both elements that are compared. To relate F elements that are not equal in all frees and bounds, we introduce a new map-restricted relator mr_rel_F . It takes a function for each free and bound - with the appropriate restrictions to small-support and bijectivity - in addition to the relations for the lives. It then uses the graphs Grp of these functions as relations for the respective free or bound variable. Transferring the ideas of bijectivity and small-support to these graph relations, the graph of a bijective function relates each atom to exactly one other atom, while the graph of small-support function acts as equality on all the arguments that are not in its support. The new arguments of the map-restricted relator are placed in front of the relations for the live variables. It is then defined in terms of the relator as shown below. Note, that relating two elements with the graph of a function v is equivalent to mapping v over the first element and relating that to the second one by equality. Thus, we define it as follows:

$$\text{mr_rel}_F u^b v^{fr} R^l x y = \text{rel}_F R^l (\text{map}_F \text{id}^l v^{fr} u^b x) y$$

2.3.1 MRBNF axioms

MRBNFs require the same axioms as BNFs with slight modifications. We take the formalized axioms from Figure 2.1 as a base and explain the differences.

For the `MAP_COMP`, `MAP_CONG` and `SET_MAP` axioms, we add the assumptions that the functions that correspond to frees and bounds are small-support functions and that the ones corresponding to bounds are additionally bijections. It means that $\text{small_supp } v^{fr} \wedge \text{small_supp } u^b \wedge \text{bijective } u^b$ is added as assumptions to these axioms.

Furthermore, while `REL_COMPP` stays unchanged, using the original relator, `IN_REL` is changed to be defined in terms of the map-restricted relator mr_rel_F as follows:

$$\begin{aligned} \text{mr_rel}_F u^b v^{fr} R^l x y &= \exists z. (\forall i. \text{set}_{F,i} z \subseteq \{(a, b). R_i a b\}) \wedge \\ &\quad \text{map}_F \text{fst}^l \text{id}^{fr} \text{id}^b z = x \wedge \text{map}_F \text{snd}^l v^{fr} u^b z = y \end{aligned}$$

2.3.2 Datatypes with bindings

MRBNFs can be used in a **binder_datatype** command to produce a datatype with bindings.

In the resulting MRBNF the free and bound type variables are required to be *large* and *regular*. Largeness is necessary to ensure that there are always fresh elements available for renaming. It is defined as the cardinality of the type being at least the bound of $F \text{ bd}_F$ for datatypes or the cardinal successor cardSucbd_F for codatatypes. In Isabelle the requirements of largeness and regularity are combined in dedicated type classes. For the cases where $\text{bd}_F = \aleph_0$, predefined `var` and `covar` implement the appropriate requirements.

We call MRBNFs declared through this command *binder-datatypes*

TODO: more + cite [1]

3 Linearizing MRBNFs

3.1 Linearizing MRBNFs

In this section we define the linearization of a MRBNF F on a subset of its *live* variables. Linearization means, that the resulting type only contains elements for which all atoms of the linearized variables are distinct. We say F is *non-repetitive* on these variables. This type is also a MRBNF with the same variable types (*live*, *free*, *bound*, *dead*), except for the linearized variables that change their type from *live* to *bound*. This means that the new map function is restricted to only allow bijective and small-support functions on these variables.

We formalize the idea of distinctness of atoms as *non-repetitiveness* on the linearized variables. In our notation we use $\text{lin} \leq l$ to refer to the number of linearized variables. Furthermore, we assume the variables that we linearize on to be the the *last lin* of the live variables. Consequently, the first $l' = l - \text{lin}$ lives of F are not linearized.

3.1.1 Non-repetitiveness

At the core of linearization lies the notion of *non-repetitiveness*. An element x of a type is considered to be non-repetitive with respect to a type variable α if it does not contain repeating α -atoms. For example, an α list is non-repetitive, if all of its α -elements it contains are pairwise distinct. To define non-repetitiveness for an arbitrary MRBNF, we express this property in terms of its map, set and relator functions. Considering α lists again, we can show a list xs to be distinct, iff for each other list ys of the same length, we can find a function f such that $ys = \text{map}_{\text{list}} f xs$. If xs were not distinct, there must exist two indices with the same α element in xs . Furthermore, there exists a ys that has different elements at these two indices and thus a function mapping xs to ys cannot exist, since it would have to map two same elements in xs to two differing ones in ys .

In Subsection 2.1.4 we proposed the idea to think about elements of a BNF (or MRBNF) F as containers with a certain shape with atoms in slots specified by the shape. Using this model, we can generalize the notion of lists having the same length to F elements having the same shape. We can express this through the relator by using the \top relation that relates every pair of arguments with each other. Thus, we give the

definition of equivalent shape and non-repetitiveness for list:

$$\begin{aligned} \text{eq_shape}_{\text{list}} x y &= \text{rel}_{\text{list}} \top x y \\ \text{nonrep}_{\text{list}} x &= \forall y. \text{eq_shape}_{\text{list}} x y \implies (\exists f. y = \text{map}_{\text{list}} f x) \end{aligned}$$

Note that we use the regular relator that only relates live variables with given relations while it requires equality for all frees and bounds.

Based on this, x is a non-repetitive element, if for all other elements y with equal shape, a function exists through which x can be mapped to y . In our example of list, this holds for all lists with distinct elements (given a second list y of same length, one can easily define a function mapping the distinct atoms of x to those of y). It does not hold for lists with repeating elements, because no f exists that could map two equal elements at different positions in this list to distinct elements in an arbitrary second list.

More interesting is the case of (α, β) alist which we only want to be non-repetitive on α . For our purpose of defining non-repetitiveness on a subset of the live variables, we fix the other live variables to be equal when defining equivalent shape. For MRBNFs with more than one live variable, we can give a definitions of *non-repetitiveness* and having *equal shape* on the last lin live variables. In that case, we consider x and y of type F to have equal shape with respect to the variables $\alpha_{l'+1} \dots \alpha_{\text{lin}}$, iff they are *equal* in the atoms corresponding to the non linearized lives and are related with \top in on the linearized variables. Consequently for the map in the nonrep definition, the id function is applied to the non linearized lives, since they are already required to be equal.

$$\begin{aligned} (\text{EQ_SHAPE}) \quad \text{eq_shape}_F^{\text{lin}} x y &= \text{rel}_F \langle (=)^{l'} (\top)^{\text{lin}} \rangle x y \\ (\text{NONREP}) \quad \text{nonrep}_F^{\text{lin}} x &= \forall y. \text{eq_shape}_F^{\text{lin}} x y \implies (\exists f^{\text{lin}}. y = \text{map}_F \langle \text{id}^{l'} f^{\text{lin}} \rangle \text{id}^{\text{fr}} \text{id}^b x) \end{aligned}$$

Note, that we use $\langle \dots \rangle$ to indicate arguments that belong together, e.g., that they are both related to lives in this case. They are just inserted to improve readability. Once again, our arguments hold for any ordering of type variables. The assumption that we linearize on the last lin variables only serves readability and is not a limitation in the actual implementation, where we allow an arbitrary subset of lives to be chosen for linearization.

3.1.2 Conditions for linearization

Two properties are necessary to linearize a MRBNFs. First, to ensure that the resulting type constructor is non-empty, it is required, that there exists a non-repetitive element (with respect to the linearized variables): $\exists x. \text{nonrep}_F^{\text{lin}} x$.

Furthermore, even though MRBNFs are already required to preserve weak pullbacks, as described in Subsection 2.1.5, for the linearization it is required that they preserve

all pullbacks. Formalized this means that the existence of z in the IN_REL axiom (Fig. 2.1) has to be fulfilled uniquely, i.e., for each R -related x and y there exists *exactly one* z fulfilling this property. For example, strong pullback preservation is fulfilled by the α list and $\alpha \beta$ prod functor but not by α fset, the type constructor for finite sets of α s.

$$\begin{aligned} (\text{PB_STRONG}) \quad \text{rel}_F R^l x y = \exists! z. & (\forall i. \text{set}_{F,i} z \subseteq \{(a, b). R_i a b\}) \wedge \\ & \text{map}_F \text{fst}^l \text{id}^{fr} \text{id}^b z = x \wedge \text{map}_F \text{snd}^l \text{id}^{fr} \text{id}^b z = y \end{aligned}$$

We note here that the requirement of strong pullback preservation can be omitted, when the MRBNF is linearized on all its live variables, i.e., when the linearized MRBNF has no live variables. This is because in this case the REL_EXCHANGE lemma explained in Subsection 3.1.3 becomes trivial. In all other cases, that lemma is the sole reason, strong pullback preservation is required.

3.1.3 Intermediate lemmas

We want to prove the MRBNF axioms for the linearized MRBNF. For this we utilize a few intermediate lemmas which we present in this section.

F is strong

From the pullback preservation with uniqueness we can prove the following lemma. In fact this notion of strength is equivalent to pullback preservation:

$$(\text{F_STRONG}) \quad \text{rel}_F R^l x y \wedge \text{rel}_F Q^l x y \implies \text{rel}_F (\inf R Q)^l x y$$

Here the infimum \inf of two relations R_i and Q_i relates exactly those elements that are related by both R_i and Q_i . To prove this lemma we first conclude that since x and y are related through some relations, they are also related with the \top -relation on all lives. This is the case since the relator or a MRBNF is monotonic and \top relates all atoms with each other. Unfolding PB_STRONG on that, we can eliminate the subset conditions for the setters, since the right sides of the subsets are just the universe of pairs with appropriate type. The reason for this is the \top -relation that is fulfilled by every pair in this set. This now gives us the knowledge that there exists exactly one z that is the "zipped" version of x and y , or in other words any two z and z' fulfilling this condition

must be equal.

$$\begin{aligned}
 & \text{rel}_F R^l x y \\
 \implies & \text{rel}_F (\top)^l x y && \text{rel}_F \text{ mono} \\
 \implies & \exists! z. (\forall i. \text{set}_{F,i} z \subseteq \{(a, b). (\top) a b\}) \wedge \\
 & \quad \text{map}_F \text{ fst}^l \text{id}^{fr} \text{id}^b z = x \wedge \text{map}_F \text{ snd}^l \text{id}^{fr} \text{id}^b z = y && \text{PB_STRONG} \\
 \implies & \exists! z. \text{map}_F \text{ fst}^l \text{id}^{fr} \text{id}^b z = x \wedge \text{map}_F \text{ snd}^l \text{id}^{fr} \text{id}^b z = y && \top \equiv \text{True} \\
 \implies & \forall z z'. (\text{map}_F \text{ fst}^l \text{id}^{fr} \text{id}^b z = x \wedge \text{map}_F \text{ snd}^l \text{id}^{fr} \text{id}^b z = y \wedge \\
 & \quad \text{map}_F \text{ fst}^l \text{id}^{fr} \text{id}^b z' = x \wedge \text{map}_F \text{ snd}^l \text{id}^{fr} \text{id}^b z' = y \implies \exists! \text{ alternative} \\
 & \quad z = z')
 \end{aligned}$$

When we unfold the MRBNF axiom `IN_REL` on the original formulation of the lemma, we obtain a construct with two existential quantifiers $\exists z$ in the assumptions and one in the goal. With the knowledge we gained above, we can conclude that they must all be equal

$$\begin{aligned}
 & \text{rel}_F R^l x y \wedge \text{rel}_F Q^l x y \implies \text{rel}_F (\inf R Q)^l x y \\
 \equiv & \exists z_R. (\forall i. \text{set}_{F,i} z_R \subseteq \{(a, b). R_i a b\}) \wedge && \text{IN_REL} \\
 & \quad \text{map}_F \text{ fst}^l \text{id}^{fr} \text{id}^b z_R = x \wedge \text{map}_F \text{ snd}^l \text{id}^{fr} \text{id}^b z_R = y \wedge \\
 \exists & z_Q. (\forall i. \text{set}_{F,i} z_Q \subseteq \{(a, b). Q_i a b\}) \wedge \\
 & \quad \text{map}_F \text{ fst}^l \text{id}^{fr} \text{id}^b z_Q = x \wedge \text{map}_F \text{ snd}^l \text{id}^{fr} \text{id}^b z_Q = y \implies \\
 \exists & z_{\inf}. (\forall i. \text{set}_{F,i} z_{\inf} \subseteq \{(a, b). (\inf R_i Q_i) a b\}) \wedge \\
 & \quad \text{map}_F \text{ fst}^l \text{id}^{fr} \text{id}^b z_{\inf} = x \wedge \text{map}_F \text{ snd}^l \text{id}^{fr} \text{id}^b z_{\inf} = y \\
 \equiv & (\forall i. \text{set}_{F,i} z \subseteq \{(a, b). R_i a b\}) \wedge && z = z' \text{ (see above)} \\
 & (\forall i. \text{set}_{F,i} z \subseteq \{(a, b). Q_i a b\}) \implies \\
 & (\forall i. \text{set}_{F,i} z \subseteq \{(a, b). (\inf R_i Q_i) a b\})
 \end{aligned}$$

The last step is proven by applying common rules on sets of pairs, subsets and conjunction after $(\inf R_i Q_i) a b$ is unfolded to $R_i a b \wedge Q_i a b$.

Relation exchange

The *exchange of relations* is a consequence of the previous property, `F_STRONG`: If two elements x and y are related through the relator with two different lists $R^l = R_1 \dots R_l$ and $Q^l = Q_1 \dots Q_l$ of atom-level relations, then x and y are also related with any

index-wise combination of R^l or Q^l . For each index i either the relation R_i or Q_i is selected.

For our purpose of linearization, we are specifically interested in the case, where for all live variables that we linearize on the relation from R^l is chosen and for all others the relation from Q^l relation, i.e., $Q^{l'} R^{\text{lin}}$. This results in the following lemma for a MRBNF F :

$$(\text{REL_EXCHANGE}) \quad \text{rel}_F R^l x y \wedge \text{rel}_F Q^l x y \implies \text{rel}_F \langle Q^{l'} R^{\text{lin}} \rangle x y$$

The lemma F_STRONG states that for each type variable the atoms are related with the infimum of R_i and Q_i . This means, that the atoms of each type variable can be related with R_i and Q_i at the same time. To prove this lemma we choose the appropriate relation from each of the l infima. This informal idea is the core of this lemma's formal proof.

In the specific case, that the MRBNF is linearized on *all* of it's live variables, $l' = 0$ and $\text{lin} = l$ resulting in R^l as the combination that is chosen. Then the lemma becomes trivial, since its goal is equal to it's first assumption in this case.

As a consequence of this, the previous lemma F_STRONG is not needed to prove this lemma. Furthermore, this lemma is the sole reason why F_STRONG and strong pullback preservation are needed for the linearization. Thus the requirement of pullback preservation can be lifted, in the case that the linearization is applied to all live variables at the same time.

Mapper preserves non-repetitiveness

An important lemma used frequently in the following proofs is that the mapper preserves non-repetitiveness. It means that given a non-repetitive element x , the result of mapping functions over it is also non-repetitive. These functions must fulfill the appropriate restrictions of bijectivity and small-support for the bounds and lives. Additionally the functions f^{lin} on the linearized lives need to be bijective.

$$(\text{NONREP_MAP}) \quad \text{small_supp } v^{fr} \wedge \text{small_supp } u^b \wedge \text{bijective } u^b \wedge \\ \text{bijective } f^{\text{lin}} \wedge \text{nonrep}_F^{\text{lin}} x \implies \text{nonrep}_F^{\text{lin}} (\text{map}_F \langle g^{l'} f^{\text{lin}} \rangle v^{fr} u^b x)$$

To give proof sketch for this lemma, we split it into two parts. First, we argue that mapping bijective f^{lin} over the linearized variables and id over all others preserves non-repetitiveness. $\text{eq_shape}_F^{\text{lin}}$ is transitive (both $=$ and \top are transitive) and x and $\text{map}_F \langle \text{id}^{l'} f^{\text{lin}} \rangle \text{id}^{fr} \text{id}^b x$ have the same shape. Thus, we have to think about the same $\forall y$ in the NONREP definition to show non-repetitiveness for the mapped x . This means,

we can fix the y and show the following, where we rename the f^{lin} in the existential quantifier to f_{E1}^{lin} and f_{E2}^{lin} to avoid naming clashes and for clarity.

$$\begin{aligned}\exists f_{E1}^{\text{lin}}. y &= \text{map}_F \langle \text{id}^{l'} f_{E1}^{\text{lin}} \rangle \text{id}^{fr} \text{id}^b x \implies \\ \exists f_{E2}^{\text{lin}}. y &= \text{map}_F \langle \text{id}^{l'} f_{E2}^{\text{lin}} \rangle \text{id}^{fr} \text{id}^b (\text{map}_F \langle \text{id}^{l'} f^{\text{lin}} \rangle \text{id}^{fr} \text{id}^b x)\end{aligned}$$

From this we obtain and fix h^{lin} from the existential quantifier $\exists f_{E1}^{\text{lin}}$ in the assumption and instantiate the existential qualifier in the goal with $f_{E2}^{\text{lin}} = (h \circ (\text{inv } f))^{\text{lin}}$. Since all f^{lin} are bijective, we know that the inverse inv for each of them exist. Using **MAP_COMP**, we can transform the instantiated goal to the assumption with the fixed h^{lin} as follows, which proves this part of the lemma:

$$\begin{aligned}y &= \text{map}_F \langle \text{id}^{l'} (h \circ (\text{inv } f))^{\text{lin}} \rangle \text{id}^{fr} \text{id}^b (\text{map}_F \langle \text{id}^{l'} f^{\text{lin}} \rangle \text{id}^{fr} \text{id}^b x) \\ &\equiv y = \text{map}_F \langle \text{id}^{l'} ((h \circ (\text{inv } f)) \circ f)^{\text{lin}} \rangle \text{id}^{fr} \text{id}^b x && \text{MAP_COMP} \\ &\equiv y = \text{map}_F \langle \text{id}^{l'} (h \circ ((\text{inv } f) \circ f))^{\text{lin}} \rangle \text{id}^{fr} \text{id}^b x && \circ \text{ assoc} \\ &\equiv y = \text{map}_F \langle \text{id}^{l'} (h \circ \text{id})^{\text{lin}} \rangle \text{id}^{fr} \text{id}^b x && \text{inv } \circ \\ &\equiv y = \text{map}_F \langle \text{id}^{l'} h^{\text{lin}} \rangle \text{id}^{fr} \text{id}^b x && \circ \text{ id}\end{aligned}$$




Figure 3.1: Illustration of lemma **NONREP_MAP**

It remains to show that mapping appropriately restricted functions over the other, not linearized type variables also preserves non-repetitiveness. From the definition **NONREP** we know that for any element y of equivalent shape, we can find functions f^{lin} acting only on the linearized variables to map to this other element. Next, we consider the mapped $x' = \text{map}_F \langle g^{l'} \text{id}^{\text{lin}} \rangle v^{fr} u^b x$ and all y' of equivalent shape to it. To proceed we show, that for every fixed y' there exists a y of equivalent shape to the original x , such that y' can be expressed as $y' = \text{map}_F \langle g^{l'} \text{id}^{\text{lin}} \rangle v^{fr} u^b y$. Knowing this, we deduce that the same f^{lin} that map x to y from the assumption also map x' to y' .

Using `MAP_COMP` this proofs non-repetitiveness for x' . Figure 3.1 shows an illustration of this lemma. The thick or bolded elements are what we know from the assumption. The dashed arrow stands for the existence of a y with equivalent shape to x for each y' on which this proof depends. This existence can be shown using `IN_REL` and other theorems about the relator.

Mapper reflects non-repetitiveness

Lastly we need the reflection of non-repetitiveness through map. This lemma can be seen as a partial reverse of `NONREP_MAP`, but only on the non-linearized lives.

$$(\text{NONREP_MAP_REV}) \quad \text{nonrep}_F^{\text{lin}} (\text{map}_F \langle g^{l'} \text{id}^{\text{lin}} \rangle \text{id}^{\text{fr}} \text{id}^b x) \implies \text{nonrep}_F^{\text{lin}} x$$

We start of the proof by defining $x' = \text{map}_F \langle g^{l'} \text{id}^{\text{lin}} \rangle \text{id}^{\text{fr}} \text{id}^b x$ and fixing a y with equivalent shape to x : $\text{eq_shape}_F^{\text{lin}} x y$. This time we can easily obtain and fix $y' = \text{map}_F \langle g^{l'} \text{id}^{\text{lin}} \rangle \text{id}^{\text{fr}} \text{id}^b y$ with $\text{eq_shape}_F^{\text{lin}} x' y'$ and $y' = \text{map}_F \langle \text{id}^{l'} f^{\text{lin}} \rangle \text{id}^{\text{fr}} \text{id}^b y$ from the assumption. We can express these maps in terms of the relator using graphs Grp of functions and *converse* graphs Grp^{-1} . For this we use the following properties obtained from the MRBNF F :

$$\begin{aligned} (\text{REL_MAP_1}) \quad & \text{rel}_F (\text{Grp } f)^l x y \equiv \text{map}_F f^l \text{id}^{\text{fr}} \text{id}^b x = y \\ (\text{REL_MAP_2}) \quad & \text{rel}_F (\text{Grp}^{-1} f)^l x y \equiv x = \text{map}_F f^l \text{id}^{\text{fr}} \text{id}^b y \end{aligned}$$

With this we can express the current proof state as shown in Figure 3.2. The arrows denote a relation from the element at their base to the element at their tip through the given relator. The top relation is obtained from $\text{eq_shape}_F^{\text{lin}} x y$, the bottom arrow with f^{lin} from the non-repetitiveness in the assumption, while the vertical arrows denote the definitions of ' x ' and ' y ' in terms of the relator.




Figure 3.2: Illustration of lemma `NONREP_MAP_REV`

Following the arrows and applying relational composition with the associated axiom `REL_COMPP` we can express the relationship path between x and y through x' and y' as $\text{rel}_F \langle (\text{Grp } g \bullet \text{Grp}^{-1} g)^{l'} (\text{Grp } f)^{\tilde{l}n} \rangle x y$. Together with the top relation from the definition of equivalent shape we can apply the `REL_EXCHANGE` lemma to receive $\text{rel}_F \langle (=)^{l'} (\text{Grp } f)^{\tilde{l}n} \rangle x y$. Unfolding the `REL_MAP_1` equality, we can show the non-repetitiveness of x :

$$\begin{aligned} & \text{rel}_F \langle (\text{Grp } g \bullet \text{Grp}^{-1} g)^{l'} (\text{Grp } f)^{\tilde{l}n} \rangle x y \wedge \text{rel}_F \langle (=)^{l'} (\top)^{\tilde{l}n} \rangle x y \\ & \equiv \text{rel}_F \langle (=)^{l'} (\text{Grp } f)^{\tilde{l}n} \rangle x y && \text{REL_EXCHANGE} \\ & \equiv y = \text{map}_F \langle \text{id}^{l'} f^{\tilde{l}n} \rangle \text{id}^{fr} \text{id}^b x && \text{REL_MAP_1} \end{aligned}$$

3.1.4 Defining the subtype and its constants

Using our definition of non-repetitiveness, we carve out a subtype of F using Isabelle's `typedef` command. This subtype F' contains exactly those elements from F that are non-repetitive on the linearized variables $\alpha_{l'+1} \dots \alpha_{\tilde{l}n}$. It furthermore provides us with the morphisms $\text{rep}_{F'}$ to convert F' elements to the type F and $\text{abs}_{F'}$ to convert F elements to F' - provided that they are non-repetitive.

In the following we specify the MRBNF constants, i.e., the mapper, setters, bound and relator for F' . We define these in terms of the base type's constants and apply the morphisms to match the types: For the setters, $\text{rep}_{F'}$ is applied to the argument before applying the appropriate setter of F to it. Since this does not change the set that is outputted, we can use the bound of F for F' . For the relator, the relations for the linearized lives are fixed to the equality relation, since in the new MRBNF these will be bounds. Lastly, for the mapper we only allow it to map bijective functions on the linearized variables in addition to the restrictions for the existing frees and bounds. This restriction is necessary to ensure that applying the map function to a F' element preserves its non-repetitiveness. If a function that violates any of the restrictions is given to the mapper, it is ignored and not applied.

As for the morphisms, concretely, we apply $\text{rep}_{F'}$ to the F' arguments of the new mapper, setters and relator, and $\text{abs}_{F'}$ to the result of the mapper. This leads us to the following definitions:

$$\begin{aligned} \text{bd}_{F'} &= \text{bd}_F \\ \text{set}_{F',i} &= \text{set}_{F,i} \circ \text{rep}_{F'} \\ \text{map}_{F'} \langle f^{l'} g^{\tilde{l}n} \rangle u^{fr} v^b &= \text{abs}_{F'} \circ (\text{map}_F \langle f^{l'} (\text{asBij } g)^{\tilde{l}n} \rangle (\text{asSS } v)^{fr} (\text{asBij } (\text{asSS } u))^b) \circ \text{rep}_{F'} \\ \text{rel}_{F'} R^{l'} x y &= \text{rel}_F \langle R^{l'} (=)^{\tilde{l}n} \rangle (\text{rep}_{F'} x) (\text{rep}_{F'} y) \end{aligned}$$

where we enforce bijectivity of the g^{bi} and u^b using $\text{asBij } f = \text{if bijective } f \text{ then } f \text{ else id}$. Analogously, both v^{fr} and u^b are enforced to be small-support functions using an analogously defined asSS .

3.1.5 Proving the MRBNF axioms

To show that F' is a MRBNF, we have to prove the axioms from Figure 2.1 for it. For most of the axioms this is straight-forward, as they only require unfolding the definitions of the new F' constants, applying the axioms of the original F and a few simple transformations. The axioms `MAP_ID`, `MAP_CONG` and `SET_BD` are proven this way, while `MAP_COMP` and `SET_MAP` require just a little more effort. Both contain the composition of $\text{map}_{F'}$ or $\text{set}_{F'}$ with map_F , respectively.

As an example we show `SET_MAP` for F' below. Note that we assume i to be in the range $1 \leq i \leq \text{ws}$ where ws is the number of all non-dead type variables, i.e., $\text{ws} = l + fr + b$. The proof works the same for setters of frees and bounds. Furthermore we assume all functions f^{ws} fulfilling their respective requirements (bijectivity and small-support) and thus all asBij and asSS evaluating to the then case.

$$\begin{aligned}
& \text{set}_{F',i}(\text{map}_{F'} f^{\text{ws}} x) \\
& \equiv \text{set}_{F,i} \circ \text{rep}_{F'}((\text{abs}_{F'} \circ (\text{map}_F f^{\text{ws}}) \circ \text{rep}_{F'}) x) && \text{unfold defs} \\
& \equiv \text{set}_{F,i}(\text{rep}_{F'}(\text{abs}_{F'}(\text{map}_F f^{\text{ws}}(\text{rep}_{F'} x)))) && \circ \text{application} \\
& \equiv \text{nonrep}_F^{\text{bi}}(\text{map}_F f^{\text{ws}}(\text{rep}_{F'} x)) \implies \text{set}_{F,i}(\text{map}_F f^{\text{ws}}(\text{rep}_{F'} x)) && \text{abs inverse} \\
& \equiv \text{nonrep}_F^{\text{bi}}(\text{rep}_{F'} x) \implies \text{set}_{F,i}(\text{map}_F f^{\text{ws}}(\text{rep}_{F'} x)) && \text{NONREP_MAP} \\
& \equiv \text{set}_{F,i}(\text{map}_F f^{\text{ws}}(\text{rep}_{F'} x)) && \text{nonrep } \text{rep}_{F'} \\
& \equiv f_i \cdot \text{set}_{F,i}(\text{rep}_{F'} x) && \text{SET_MAP of } F \\
& \equiv f_i \cdot \text{set}_{F',i} x && \text{fold defs, } \circ
\end{aligned}$$

where "abs inverse" denotes the theorem that $\text{rep}_{F'}$ is the inverse of $\text{abs}_{F'}$ for arguments that are non-repetitive. Furthermore "nonrep $\text{rep}_{F'}$ " states that converting a F' element to F inherently means that the F element is non-repetitive.

The validity of the bound `BD` is trivially proven, since the bound is copied from F .

It remains to show `REL_COMPP` and `IN_REL` for F' . While the former is easily proven using the corresponding axiom of F and some simple properties of relational composition, the latter is certainly the most interesting axiom to show.

We do not show a full proof of this property here, but investigate an interesting step. In the proof we reach a state, where we need to show that $\text{nonrep}_F^{\text{bi}}(\text{map}_F \text{fst}^l \text{id}^{\text{fr}} \text{id}^b z) \implies \text{nonrep}_F^{\text{bi}}(\text{map}_F \langle \text{id}^l \text{fst}^{\text{bi}} \rangle \text{id}^{\text{fr}} \text{id}^b z)$. To give an intuition for why this is necessary,

we obtain the left side of the implication from the `IN_REL` axiom of F and need to show the right side to eliminate a composition $\text{abs}_{F'} \circ \text{rep}_{F'}$ in the goal state.

The step is proven as follows:

$$\begin{aligned} \text{nonrep}_F^{\text{lin}} (\text{map}_F \text{ fst}^l \text{ id}^{fr} \text{ id}^b z) &\implies \\ \text{nonrep}_F^{\text{lin}} (\text{map}_F \langle \text{fst}^l \text{ id}^{\text{lin}} \rangle \text{ id}^{fr} \text{ id}^b (\text{map}_F \langle \text{id}^l \text{ fst}^{\text{lin}} \rangle \text{ id}^{fr} \text{ id}^b z)) &\implies \\ \text{nonrep}_F^{\text{lin}} (\text{map}_F \langle \text{id}^l \text{ fst}^{\text{lin}} \rangle \text{ id}^{fr} \text{ id}^b z) \end{aligned}$$

The first step is reached through `MAP_COMP` of F , while the second one needs the `NONREP_MAP_REV` lemma. This is the final place, where strong pullback preservation is used and the reason why it is required.

3.1.6 Lifting Witnesses

Existing witnesses of the original MRBNF that do not depend on any of the linearized variables can be lifted to be witnesses of the linearized MRBNF.

For this it is necessary to show that they are non-repetitive on the linearized elements, i.e., that they are part of the new type. From `WITS` (Subsection 2.1.6) we know that any witness not depending on the linearized lives does not contain atoms from these lives. Thus, we can show that these witnesses are non-repetitive, since an element with no α atoms is trivially non-repetitive on α .

Other witnesses that depend on the linearized variables cannot be lifted and have to be discarded. Even if they are non-repetitive, witnesses of a MRBNF may only depend on lives and not on bounds, which the linearized lives turn into.

Additionally, new witnesses may be specified for the resulting MRBNF. For these the property `WITS` defined in Subsection 2.1.6 has to be proven, i.e., that they only consist of the atoms given to them as arguments. Furthermore, it has to be shown that they are part of the type, i.e., that they are non-repetitive.

When an liftable witness of the original MRBNF exists or a new witness fulfilling `WITS` is specified, the existence of a non-repetitive element we motivated in Subsection 3.1.2 is trivially proven.

3.1.7 Preservation of strength

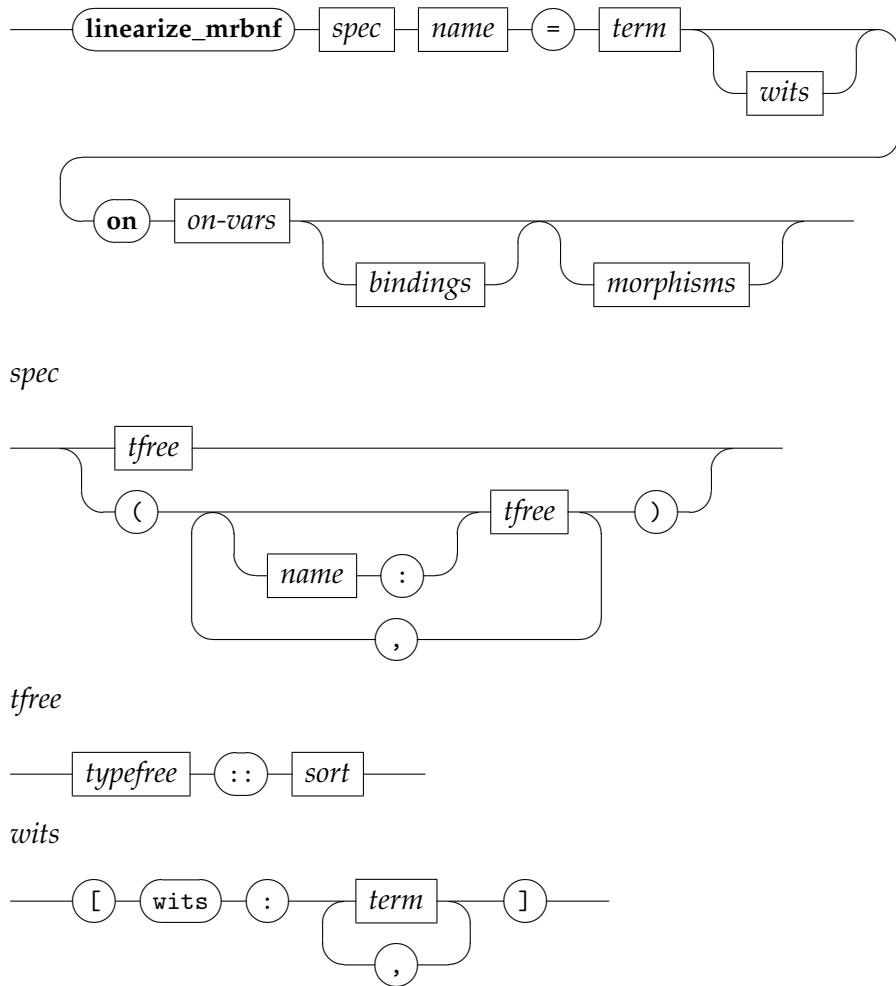
Linearizing an MRBNF preserves the strength property of it. This means, that for the new F' the following axiom holds, provided that F fulfills `F_STRONG`:

$$\text{rel}_{F'} R'^l x y \wedge \text{rel}_{F'} Q'^l x y \implies \text{rel}_{F'} (\inf R Q)^l x ys$$

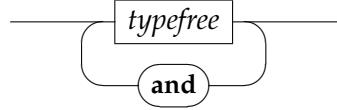
This is easily proven by unfolding the definition of rel_F , applying F_STRONG and unfolding $(\text{inf } (=) (=)) \equiv (=)$ for the linearized variables. Strength of an MRBNF is a property that often comes in useful, however in Isabelle it is not tracked in the MRBNF construct at the moment. At the very least, this proof allows us to easily linearize a linearized MRBNF again on further live variables.

3.2 Implementing the linearize_mrbnf command

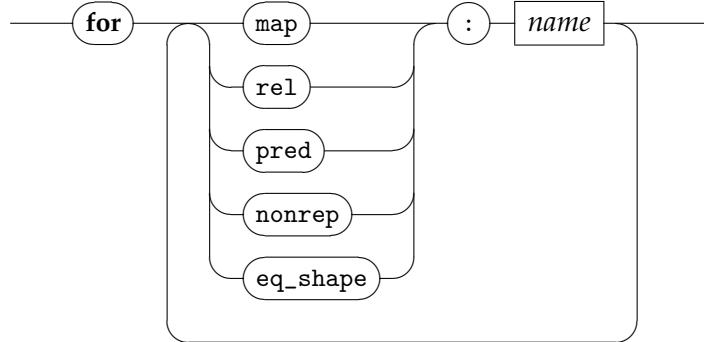
We implement a command that allows the user to linearize an existing MRBNF or BNF on one or multiple of its live variables. The syntax of the command is given in the following:



on-vars



bindings



morphisms



With this command, we can linearize our example by writing the following line in Isabelle:

`linearize_mrbnf (keys: $\alpha :: \text{var}$, vals: β) alist = " $(\alpha :: \text{var} \times \beta)$ list" on α`

Since for $(\alpha \times \beta)$ list both type variables are live and we only linearize on α , it is necessary to prove strong pullback preservation for this MRBNF.

After the user has written the command, the conditions for linearization we presented in Subsection 3.1.2 have to be proven. These are the non-emptiness of the linear type, strong pullback preservation PB_STRONG and the witness axioms WITS if applicable.

These conditions dynamically generated for the user. For example, it is only necessary to show strong pullback preservation PB_STRONG, when the resulting MRBNF has live variables remaining. Furthermore, as mentioned in Subsection 3.1.6, the non-emptiness of the non-repetitive type is easily proven when the user specified a non-emptiness witness, or a liftable witness of the original type exists. The user is not asked to show the goals that are actually necessary for a specific linearization.

Furthermore, since the original MRBNF already fulfills weak pullback preservation, we extract the uniqueness property of strong pullback preservation and require the

user to prove only this. Strong pullback preservation PB_STRONG can be proven from weak pullback preservation IN_REL together with the uniqueness property we specify as follows:

$$(PB_UNIQUE) \quad \forall x y. (\text{map}_F \text{ fst}^l \text{ id}^{fr} \text{ id}^b x = \text{map}_F \text{ fst}^l \text{ id}^{fr} \text{ id}^b y \wedge \\ \text{map}_F \text{ snd}^l \text{ id}^{fr} \text{ id}^b x = \text{map}_F \text{ snd}^l \text{ id}^{fr} \text{ id}^b y) \implies \\ x = y$$

When the required properties are shown, the subtype and the constants for the new MRBNF are defined.

The proofs of the intermediate lemmas from Subsection 3.1.3 and MRBNF axioms are automated through ML-tactics. For this the existing high-level apply-style and Isar proofs were converted to single-step apply proofs. These proofs avoid using the automatic proof tactics of Isabelle like `metis`, `auto`, `fastforce` and even `simp`. Instead they rely on explicit rule applications, substitutions, deterministic repetitions and in certain cases instantiations of free variables in existing theorems and lemmas.

4 Examples

4.1 POPLmark Challenge: Pattern

The POPLmark Challenge [6] presents a selection of problems to benchmark the progress in formalizing programming language metatheory. The challenges are built around formalizing aspects of *System F*_< calculus, a polymorphic typed lambda calculus with subtyping. We are interested in part 2B of this Challenge, which has the goal to formalize and proof *type soundness* for terms with pattern matching over records. Type soundness is considered in terms of *preservation* (evaluating a term preserves its type) and *progress* (a term is either a value or can be evaluated).

The challenge defines the syntax of terms t and patterns p as follows:

(variable)	$t ::= x$
(abstraction)	$ \lambda x : T. t$
(application)	$ t_1 t_2$
(type abstraction)	$ \lambda X <: T. t$
(type application)	$ t [T]$
(record)	$ \{j \in 1 \dots n : l_j = t_j\}$
(projection)	$ t.l$
(pattern-let)	$ \text{let } p = t_1 \text{ in } t_2$
(variable pattern)	$p ::= x : T$
(record pattern)	$ \{j \in 1 \dots n : l_j = p_j\}$

In this syntax types T are defined similarly, however we omit this here as it is not important for this example. Furthermore x stand for variables, X for type variables and l for labels.

We focus on the *record* and *pattern-let* terms together with the *variable* and *record patterns*. A record is a term defined as a finite set of assignments of n terms to one label each. The labels l within a record must be pairwise distinct. A pattern is defined as either a typed variable or a finite set of n assignments of patterns to labels. Again the labels must be pairwise distinct.

A formalization of part 2B of the POPLmark Challenge in Isabelle/HOL is presented by Blanchette et al. [1]. Datatypes with bindings play an important role in their formalization. They represent types T as α typ and terms t as (α, β) trm. These two are defined with the **binder_datatype** command, where α is the type representing type variables X and β represents variable names x . Labels are implemented as strings, however any infinite type could be used in their place.

A central notion in this formalization is the *labeled finite set* (α, β) lfset that is used in the representation of records and record patterns. This type constructor is a subtype of $(\alpha \times \beta)$ fset that only includes elements that are non-repetitive on α . This restriction is necessary, because for both records and patterns the label α must be mutually distinct, i.e., the set representing them has to be non-repetitive.

While by construction $(\alpha \times \beta)$ fset is a BNF (and an MRBNF since all BNFs are also MRBNFs) with both variables being live, (α, β) lfset is a MRBNFs with α as a bound variable, since it is non-repetitive on α . While this is a linearization, the finite set on pairs does not fulfill strong pullback preservation. Thus the approach and command we presented in Chapter 3 cannot be used here. Because of an alternate, equivalent description on non-repetitiveness specific to this type, it is still possible to manually linearize this MRBNF. The linearized lfset is used in the **binder_datatype** definition of term trm as one the alternative representing records:

```
| Rec "(string, (\alpha, \beta) trm) lfset"
```

For the pattern a new type is used. It is constructed by linearizing an intermediate datatype prepat that is defined using the **datatype** command:

```
datatype (\alpha, \beta) prep = PPVar "\beta" "\alpha typ" | PPRRec "(string, (\alpha, \beta) prep) lfset"
```

This datatype is a MRBNF with α as a free and β as a live variable. The binder-datatype typ is a unary MRBNF with its only type variable being free, prep inherits this variable type for α .

This datatype can represent all possible patterns just fine, however it also represents many patterns that are not allowed. In the definition of Challenge 2B it is stated that “[...] the variable patterns appearing in a pattern are assumed to bind pairwise distinct variables” ([6]). Thus, we linearize (α, β) prep on the type variables for β

This MRBNF can be linearized using the **linearize_mrbnf** command from Chapter 3 as follows:

```
linearize_mrbnf (PTVars: \alpha :: var, PVars: \beta :: var) pat = "(\alpha :: var, \beta :: var) prep"
  [wits: "PPRec lfempty :: (\alpha :: var, \beta :: var) prep"] on \beta
```

We specified the non-emptiness witness "PPRec lfempty" that is constructed by applying the PPRec constructor of prepat to the empty lfset lfempty. This witness is independent of any type variable and thus it trivially fulfills the witness axiom WITS.

The command generates the conditions that need to be proven. Since β is the only live of prepat, strong pullback preservation PB_STRONG (or PB_UNIQUE) does not have to be proven as explained in Subsection 3.1.3. Furthermore, since we specified a non-emptiness witness, we do not have to proof that a non-repetitive element of the type exists, but show that the witness "PPRec lfempty" is non-repetitive instead. This is easily shown with the MRBNF-axioms after noticing that "PPRec lfempty" is only equivalent in shape to itself.

The linearized MRBNF is then used in the definition of trm in the alternative that represents pattern-lets:

| Let " $(\alpha, p::\beta)$ pat" " (α, β) trm" $t::"(\alpha, \beta)$ trm" binds p in t

Abbreviations

BNF Bounded Natural Functor

MRBNF Map-Restricted Bounded Natural Functor

List of Figures

2.1	The BNF axioms	4
2.2	$\text{set}_{F,i}$ as a natural transformation	5
3.1	Illustration of lemma NONREP_MAP	18
3.2	Illustration of lemma NONREP_MAP_REV	19

List of Tables

Bibliography

- [1] J. C. Blanchette, L. Gheri, A. Popescu, and D. Traytel, “Bindings as bounded natural functors,” *Proc. ACM Program. Lang.*, vol. 3, no. POPL, Jan. 2019. doi: 10.1145/3290335.
- [2] D. Traytel, A. Popescu, and J. C. Blanchette, “Foundational, compositional (co)datatypes for higher-order logic: Category theory applied to theorem proving,” in *2012 27th Annual IEEE Symposium on Logic in Computer Science*, 2012, pp. 596–605. doi: 10.1109/LICS.2012.75.
- [3] J. van Brügge, A. Popescu, and D. Traytel, “Animating MRBNFs: Truly Modular Binding-Aware Datatypes in Isabelle/HOL,” in *16th International Conference on Interactive Theorem Proving (ITP 2025)*, Y. Forster and C. Keller, Eds., ser. Leibniz International Proceedings in Informatics (LIPIcs), vol. 352, Dagstuhl, Germany: Schloss Dagstuhl – Leibniz-Zentrum für Informatik, 2025, 11:1–11:20, ISBN: 978-3-95977-396-6. doi: 10.4230/LIPIcs.ITP.2025.11.
- [4] J. C. Blanchette, A. Popescu, and D. Traytel, “Cardinals in isabelle/hol,” in *Interactive Theorem Proving*, G. Klein and R. Gamboa, Eds., Cham: Springer International Publishing, 2014, pp. 111–127, ISBN: 978-3-319-08970-6.
- [5] J. C. Blanchette, A. Popescu, and D. Traytel, “Witnessing (co)datatypes,” in *Programming Languages and Systems*, J. Vitek, Ed., Berlin, Heidelberg: Springer Berlin Heidelberg, 2015, pp. 359–382, ISBN: 978-3-662-46669-8.
- [6] B. E. Aydemir, A. Bohannon, M. Fairbairn, J. N. Foster, B. C. Pierce, P. Sewell, D. Vytiniotis, G. Washburn, S. Weirich, and S. Zdancewic, “Mechanized metatheory for the masses: The poplmark challenge,” in *Theorem Proving in Higher Order Logics*, J. Hurd and T. Melham, Eds., Berlin, Heidelberg: Springer Berlin Heidelberg, 2005, pp. 50–65, ISBN: 978-3-540-31820-0.